

Assignment 2 Walkthrough

Warren Mansur

Reminder: Purpose of Walkthrough

- The goal of this walkthrough is to help you successfully understand and complete Assignment 2.

Reminder: Purpose of Walkthrough

- The goal of this walkthrough is to help you successfully understand and complete Assignment 2.
- I briefly recap important concepts relevant to the assignment, then show you relevant code examples.

Reminder: Purpose of Walkthrough

- The goal of this walkthrough is to help you successfully understand and complete Assignment 2.
- I briefly recap important concepts relevant to the assignment, then show you relevant code examples.
- These code examples are not shown elsewhere (i.e. not in the instructor's live class, shared example code, or online modules).

Reminder: Live Classrooms

- It's important to attend Prof. Joner's live classrooms to learn the core concepts. The sessions are thorough and deep.

Reminder: Live Classrooms

- It's important to attend Prof. Jones's live classrooms to learn the core concepts. The sessions are thorough and deep.
- This walkthrough recaps some concepts to focus on the practical application for the assignment, but does not go into the same conceptual depth.

Reminder: Live Classrooms

- It's important to attend Prof. Joner's live classrooms to learn the core concepts. The sessions are thorough and deep.
- This walkthrough recaps some concepts to focus on the practical application for the assignment, but does not go into the same conceptual depth.
- Prof. Joner provides one sample R file per week to get you started. They demonstrate use of some important libraries and functions, but still require you to apply the code to the specific assignment scenarios.

Reminder: Submitting

- It's important to provide two files – one for your answers without the code, and the other with the code.

Reminder: Submitting

- It's important to provide two files – one for your answers without the code, and the other with the code.
- This keeps your work and answers separate, helping ensure that you receive credit where it is due.

Reminder: Submitting

- It's important to provide two files – one for your answers without the code, and the other with the code.
- This keeps your work and answers separate, helping ensure that you receive credit where it is due.
- There have been occasions in past runnings where the work over-cluttered the answers, and a few lost some points because of misidentification.

Reminder: Submitting

- It's important to provide two files – one for your answers without the code, and the other with the code.
- This keeps your work and answers separate, helping ensure that you receive credit where it is due.
- There have been occasions in past runnings where the work over-cluttered the answers, and a few lost some points because of misidentification.
- Use this YAML to hide code globally:

```
1 # ----- in the document YAML -----  
2 execute:  
3   echo: false # hide code globally  
4   warning: false # hide warnings globally
```

Objects with Only Nominal Attributes

Table 1: Is Car C1 closer to C6 or C9?

ID	Make	Body Style	Fuel	Transmission	Color	Ownership
C1	Toyota	sedan	gas	automatic	white	first-owner
C2	Ford	truck	diesel	automatic	blue	corporate
C3	Honda	hatchback	gas	manual	red	second-owner
C4	BMW	coupe	gas	automatic	black	first-owner
C5	Tesla	sedan	electric	automatic	silver	corporate
C6	Hyundai	SUV	gas	automatic	white	rental
C7	Chevrolet	truck	gas	manual	red	second-owner
C8	Toyota	SUV	hybrid	automatic	blue	first-owner
C9	Ford	coupe	gas	manual	black	corporate

Hamming Distance Between Objects

- **Objective** – Quantify how similar two records are when every attribute is **nominal**.

Hamming Distance Between Objects

- **Objective** – Quantify how similar two records are when every attribute is **nominal**.
- **Representation** – Treat as n-element string of labels (only exact matches matter).

Hamming Distance Between Objects

- **Objective** – Quantify how similar two records are when every attribute is **nominal**.
- **Representation** – Treat as n-element string of labels (only exact matches matter).
- **Scaled Hamming Distance** – $\text{distance} = \# \text{ mismatched attributes} / \# \text{ attributes}$

Hamming Distance Between Objects

- **Objective** – Quantify how similar two records are when every attribute is **nominal**.
- **Representation** – Treat as n-element string of labels (only exact matches matter).
- **Scaled Hamming Distance** – $\text{distance} = \# \text{ mismatched attributes} / \# \text{ attributes}$
- 0 = all attributes match (identical profiles)

Hamming Distance Between Objects

- **Objective** – Quantify how similar two records are when every attribute is **nominal**.
- **Representation** – Treat as n-element string of labels (only exact matches matter).
- **Scaled Hamming Distance** – $\text{distance} = \# \text{ mismatched attributes} / \# \text{ attributes}$
- 0 = all attributes match (identical profiles)
- 1 = all attributes differ (maximally dissimilar)

Distance Between Cars Code

```
1 # Load the car-ownership data and compute scaled Hamming distances
2 cars <- data.frame(
3   ID       = c("C1", "C2", "C3", "C4", "C5", "C6", "C7", "C8", "C9"),
4   Make     = c("Toyota", "Ford", "Honda", "BMW", "Tesla", "Hyundai", "Chevrolet", "Toyota", "Ford"),
5   Body_Style = c("sedan", "truck", "hatchback", "coupe", "sedan", "SUV", "truck", "SUV", "coupe"),
6   Fuel     = c("gas", "diesel", "gas", "gas", "electric", "gas", "gas", "hybrid", "gas"),
7   Transmission = c("automatic", "automatic", "manual", "automatic", "automatic", "automatic", "manual", "automatic", "manual"),
8   Color     = c("white", "blue", "red", "black", "silver", "white", "red", "blue", "black"),
9   Ownership  = c("first-owner", "corporate", "second-owner", "first-owner", "corporate", "rental", "second-owner", "first-owner", "corporate"),
10  stringsAsFactors = FALSE      # keep as plain strings for direct comparison
11 )
12
13 nominal_cols <- c("Make", "Body_Style", "Fuel", "Transmission", "Color", "Ownership") # Columns to compare (all six are nominal)
14 hamming <- function(id_a, id_b) { # Scaled Hamming distance = proportion of mismatches across the six features
15   i <- match(id_a, cars$ID)      # row number of id_a
16   j <- match(id_b, cars$ID)      # row number of id_b
17   mean(cars[i, nominal_cols] != cars[j, nominal_cols])
18 }
19
20 dist_C1_C6 <- hamming("C1", "C6") # Compare C1 to C6 and C9
21 dist_C1_C9 <- hamming("C1", "C9")
22 closer_car <- ifelse(dist_C1_C6 < dist_C1_C9, "C6", "C9") # Identify which car is closer to C1
23
24 list(distance_C1_C6 = dist_C1_C6, # Return a tidy summary
25       distance_C1_C9 = dist_C1_C9,
26       closer_to_C1   = closer_car)
```

Distance Between Cars Results

```
$distance_C1_C6
```

```
[1] 0.5
```

```
$distance_C1_C9
```

```
[1] 0.8333333
```

```
$closer_to_C1
```

```
[1] "C6"
```

Objects with Only Numeric Attributes

Table 2: Calculate Manhattan and Euclidean Distance for O1, O2

Object	Z1	Z2	Z3	Z4	Z5
O1	95	40	35	5	50
O2	82	70	16	4	49

Euclidean and Manhattan Distances Between Objects

- **Objective** – measure numeric dissimilarity between O1 and O2 across multiple attributes.

Euclidean and Manhattan Distances Between Objects

- **Objective** – measure numeric dissimilarity between O1 and O2 across multiple attributes.
- **Coordinate view** – treat each object as the point and each gap is computed component-wise.

Euclidean and Manhattan Distances Between Objects

- **Objective** – measure numeric dissimilarity between O1 and O2 across multiple attributes.
- **Coordinate view** – treat each object as the point and each gap is computed component-wise.
- **Manhattan distance** – sum of absolute gaps. Captures the cost of moving only along axes; less sensitive to one extreme attribute.

Euclidean and Manhattan Distances Between Objects

- **Objective** – measure numeric dissimilarity between O1 and O2 across multiple attributes.
- **Coordinate view** – treat each object as the point and each gap is computed component-wise.
- **Manhattan distance** – sum of absolute gaps. Captures the cost of moving only along axes; less sensitive to one extreme attribute.
- **Euclidean (straight-line) distance** – square-root of squared gaps. Squaring emphasizes large differences.

Distance Between O1, O2 Code

```
1  # Load the two-row Z-values data set and compute Euclidean & Manhattan distances
2  z <- data.frame(
3    Object = c("O1", "O2"),
4    Z1 = c(95, 82),
5    Z2 = c(40, 70),
6    Z3 = c(35, 16),
7    Z4 = c(5, 4),
8    Z5 = c(50, 49)
9  )
10
11 # Keep only the numeric attributes for distance calculations
12 z_vals <- z[, -1]      # drop the identifier column
13
14 # Euclidean distance - dist() defaults to Euclidean for numeric data frames
15 euclidean_O1_O2 <- as.numeric(dist(z_vals)[1])
16
17 # Manhattan ( l1 ) distance - specify method = "manhattan"
18 manhattan_O1_O2 <- as.numeric(dist(z_vals, method = "manhattan")[1])
19
20 # Return both distances in a tidy structure
21 list(euclidean = euclidean_O1_O2,
22      manhattan = manhattan_O1_O2)
```

Distance Between O1, O2 Results

```
$euclidean
```

```
[1] 37.84178
```

```
$manhattan
```

```
[1] 64
```

Package Delivery Dataset

Table 3: Package Delivery Dataset – 8 Row Sample

ID	Dist	WtLb	Fragil	TwoDay	PkRsh	WthrAd	DrvExp	HndOff	UrbDst	Status
P1	939.5	106.2	0	0	1	0	4.2	0	1	LateMinor
P2	2377.0	23.3	0	0	1	0	2.6	0	1	LateMinor
P3	1831.3	86.7	0	0	0	0	27.3	0	0	OnTime
P4	1498.7	91.2	0	0	1	0	6.1	1	1	OnTime
P5	394.3	63.9	0	0	1	0	24.7	3	1	OnTime
P6	394.2	110.6	0	0	1	0	16.6	1	0	OnTime
P7	149.9	140.2	0	0	1	0	2.8	2	0	LateMinor
P19	1082.7	8.7	0	0	1	0	26.7	2	1	LateMajor

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = DELIV_STATUS).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = DELIV_STATUS).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.
2. Fit a k-Nearest Neighbor model.

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = DELIV_STATUS).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.
2. Fit a k-Nearest Neighbor model.
 - Center and scale all predictors.

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = DELIV_STATUS).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.
2. Fit a k-Nearest Neighbor model.
 - Center and scale all predictors.
 - Use cross-validation to choose the optimal k .

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = `DELIV_STATUS`).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.
2. Fit a k-Nearest Neighbor model.
 - Center and scale all predictors.
 - Use cross-validation to choose the optimal k .
 - Predict the hold-out set.

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = `DELIV_STATUS`).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.
2. Fit a k-Nearest Neighbor model.
 - Center and scale all predictors.
 - Use cross-validation to choose the optimal k .
 - Predict the hold-out set.
 - Produce a confusion matrix and accuracy.

Building and Evaluating a k-Nearest Neighbor Classifier

Use the dataset **package_delivery.csv** (824 observations, 10 features; class = `DELIV_STATUS`).

1. Generate training and hold-out partitions, keeping $\frac{1}{3}$ of the data for hold-out.
2. Fit a k-Nearest Neighbor model.
 - Center and scale all predictors.
 - Use cross-validation to choose the optimal k .
 - Predict the hold-out set.
 - Produce a confusion matrix and accuracy.
3. Decide whether the resulting classifier is acceptable and explain why.

k-Nearest Neighbor Classifier

- **Objective** – classifies with a distance-based learner (k-NN).

k-Nearest Neighbor Classifier

- **Objective** – classifies with a distance-based learner (k-NN).
- **Train/hold-out split** – keeps evaluation unbiased; $\frac{1}{3}$ hold-out follows the holdout-method guideline.

k-Nearest Neighbor Classifier

- **Objective** – classifies with a distance-based learner (k-NN).
- **Train/hold-out split** – keeps evaluation unbiased; $\frac{1}{3}$ hold-out follows the holdout-method guideline.
- **Center & scale** – standardization ensures every numeric attribute contributes equally to distance

k-Nearest Neighbor Classifier

- **Objective** – classifies with a distance-based learner (k-NN).
- **Train/hold-out split** – keeps evaluation unbiased; $\frac{1}{3}$ hold-out follows the holdout-method guideline.
- **Center & scale** – standardization ensures every numeric attribute contributes equally to distance
- **Select k** – use cross-validation to find the k that maximizes validation accuracy; small k risks noise, large k risks over-smoothing.

k-Nearest Neighbor Classifier

- **Objective** – classifies with a distance-based learner (k-NN).
- **Train/hold-out split** – keeps evaluation unbiased; $\frac{1}{3}$ hold-out follows the holdout-method guideline.
- **Center & scale** – standardization ensures every numeric attribute contributes equally to distance
- **Select k** – use cross-validation to find the k that maximizes validation accuracy; small k risks noise, large k risks over-smoothing.
- **Confusion matrix & accuracy** – summarize predictions on the hold-out set; $\text{accuracy} = \frac{TP+TN}{\text{Total}}$, but also inspect class-wise recall and precision for balance.

k-Nearest Neighbor Code - Generating partitions

```
1 df <- read.csv("package_delivery.csv")
2
3 # Ensure the class column is categorical so caret treats it correctly.
4 df$DELIV_STATUS <- factor(df$DELIV_STATUS)
5
6 # — 2. Train / hold-out split (% : %) —————
7 # initial_split(): stratified on DELIV_STATUS, matching the hold-out
8 # method depicted on slide 11 of the lecture.
9 library(rsample)
10 set.seed(42) #Optional. Makes it precisely reproducible.
11 split <- initial_split(df, prop = 2/3, strata = DELIV_STATUS)
12 train <- training(split)
13 holdout <- testing(split)
```


k-Nearest Neighbor Code - Fit, Center, Scale, Best K

```
1 # caret::train(): mirrors lines ~410-420 of Module 2.R
2 library(caret)
3 ctrl <- trainControl(method = "cv",      # 10-fold CV
4                       number = 10,
5                       classProbs = FALSE,
6                       summaryFunction = defaultSummary)
7
8 knn_mod <- train(DELIV_STATUS ~ ., data = train,
9                 method      = "knn",
10                 trControl   = ctrl,
11                 preProcess  = c("center", "scale"), # mandated by lecture
12                 tuneLength  = 30)                  # search 30 odd k's
13
14 # Best k chosen by caret
15 best_k <- knn_mod$bestTune$k
```

k-Nearest Neighbor Code - Predict, Confusion Matrix, Accuracy

```
1  pred <- predict(knn_mod, newdata = holdout, type = "raw")
2
3  # confusionMatrix(): same metric block used in Module 2.R
4  cm <- confusionMatrix(data = pred,
5                        reference = holdout$DELIV_STATUS,
6                        positive = "OnTime") # pick the "positive" label as needed
7
8  list(
9    best_k      = best_k,
10   confusion_table = cm$table,
11   accuracy      = cm$overall["Accuracy"]
12 )
```

k-Nearest Neighbor Results

```
$best_k
```

```
[1] 31
```

```
$confusion_table
```

Reference

Prediction LateMajor LateMinor OnTime

LateMajor 0 0 0

LateMinor 10 88 12

OnTime 0 47 119

```
$accuracy
```

```
Accuracy
```

```
0.75
```

Predicting Rarer But Significant Class Values

- Sometimes in real data, a particular value appears rarely, but is still significant.

Predicting Rarer But Significant Class Values

- Sometimes in real data, a particular value appears rarely, but is still significant.
- For example, in the package dataset, very late packages (LateMajor) appear only 30 times, but its negative impact on the courier is significant.

Predicting Rarer But Significant Class Values

- Sometimes in real data, a particular value appears rarely, but is still significant.
- For example, in the package dataset, very late packages (LateMajor) appear only 30 times, but its negative impact on the courier is significant.
- One option for predicting these is binary logistic regression with over sampling.

Predicting Rarer But Significant Class Values

- Sometimes in real data, a particular value appears rarely, but is still significant.
- For example, in the package dataset, very late packages (LateMajor) appear only 30 times, but its negative impact on the courier is significant.
- One option for predicting these is binary logistic regression with over sampling.
- The rare value becomes the “Yes” or “1”.

Binary Logistic Regression

- Predicts a yes/no outcome, giving a probability between 0 and 1.

Binary Logistic Regression

- Predicts a yes/no outcome, giving a probability between 0 and 1.
- If the probability passes the cutoff, we label the row “yes”.

Binary Logistic Regression

- Predicts a yes/no outcome, giving a probability between 0 and 1.
- If the probability passes the cutoff, we label the row “yes”.
- Under-sampling randomly drops many “no” rows until yes/no are roughly even.

Binary Logistic Regression

- Predicts a yes/no outcome, giving a probability between 0 and 1.
- If the probability passes the cutoff, we label the row “yes”.
- Under-sampling randomly drops many “no” rows until yes/no are roughly even.
- Over-sampling copies or synthetically creates “yes” rows until row counts are even.

Applying Logistic Regression/Over-Sampling to Package Delivery

We want to answer this question: Given it's late, can we tell if it is majorly late?

1. Remove all OnTime values.

Applying Logistic Regression/Over-Sampling to Package Delivery

We want to answer this question: Given it's late, can we tell if it is majorly late?

1. Remove all OnTime values.
2. Split the dataset into training and holdout.

Applying Logistic Regression/Over-Sampling to Package Delivery

We want to answer this question: Given it's late, can we tell if it is majorly late?

1. Remove all OnTime values.
2. Split the dataset into training and holdout.
3. Apply over-sampling for the LateMajor value.

Applying Logistic Regression/Over-Sampling to Package Delivery

We want to answer this question: Given it's late, can we tell if it is majorly late?

1. Remove all OnTime values.
2. Split the dataset into training and holdout.
3. Apply over-sampling for the LateMajor value.
4. Fit a logistic regression model to the updated dataset.

Applying Logistic Regression/Over-Sampling to Package Delivery

We want to answer this question: Given it's late, can we tell if it is majorly late?

1. Remove all OnTime values.
2. Split the dataset into training and holdout.
3. Apply over-sampling for the LateMajor value.
4. Fit a logistic regression model to the updated dataset.
5. Generate a confusion matrix and compute the accuracy and the F-score to make predictions.

Majorly Late? Steps 1 - 3 Code

```
1 # 1. Read the package-delivery data and drop the OnTime rows
2 df <- read.csv("package_delivery.csv")
3 df <- subset(df, DELIV_STATUS != "OnTime") # keep only LateMinor / LateMajor
4 df$DELIV_STATUS <- factor(df$DELIV_STATUS) # be sure class is a factor
5
6 # 2. Stratified train / hold-out split (2/3 train, 1/3 test)
7 library(rsample)
8 set.seed(31)
9 split <- initial_split(df, prop = 0.67, strata = DELIV_STATUS)
10 train <- training(split)
11 hold <- testing(split)
12
13 # 3. Over-sample the minority class (LateMajor) in the training set
14 library(ROSE) # supplies ovun.sample()
15 train_over <- ovun.sample(DELIV_STATUS ~ .,
16                           data = train,
17                           method = "over", # simple random over-sampling
18                           seed = 33,
19                           p = 0.5)$data # make the two classes ≈50 / 50
20 table(train$DELIV_STATUS) # before
21 table(train_over$DELIV_STATUS) # after
22
23 # 4. Fit a logistic regression model on the over-sampled training data
24 logit_mod <- glm(DELIV_STATUS ~ ., data = train_over,
25                 family = "binomial") # LateMajor is reference level
```

Majorly Late? Steps 1 - 3 Results

LateMajor LateMinor

20 268

LateMinor LateMajor

268 273

Majorly Late? Steps 4 - 5 Code Part 1

```
1 # 4. Fit a logistic regression model on the over-sampled training data
2 logit_mod <- glm(DELIV_STATUS ~ ., data = train_over,
3                 family = "binomial")           # LateMajor is reference level
4
5 # 5. Predict the hold-out set and evaluate
6 prob <- predict(logit_mod, newdata = hold, type = "response")
7 pred <- factor(ifelse(prob >= 0.5, "LateMajor", "LateMinor"),
8               levels = levels(hold$DELIV_STATUS))
9
10 library(caret)
11 cm <- confusionMatrix(data = pred,
12                       reference = hold$DELIV_STATUS,
13                       positive = "LateMajor")    # treat LateMajor as "yes"
14 acc <- cm$overall["Accuracy"]
```

Majorly Late? Steps 4 - 5 Results Part 1

Confusion Matrix and Statistics

```

      Reference
Prediction LateMajor LateMinor
LateMajor      8      26
LateMinor      2     106
```

Accuracy : 0.8028

95% CI : (0.7278, 0.8648)

No Information Rate : 0.9296

P-Value [Acc > NIR] : 1

Kappa : 0.2859

McNemar's Test P-Value : 1.383e-05

Sensitivity : 0.80000

Specificity : 0.80303

Pos Pred Value : 0.23529

Neg Pred Value : 0.98148

Prevalence : 0.07042

Detection Rate : 0.05634

Detection Prevalence : 0.23944

Balanced Accuracy : 0.80152

'Positive' Class : LateMajor

Majorly Late? Steps 4 - 5 Code Part 2

```
1 # F-score helper (mirrors Module 2 script)
2 calc_measures <- function(cm_tbl) {
3   tp <- cm_tbl["LateMajor", "LateMajor"]
4   fp <- cm_tbl["LateMajor", "LateMinor"]
5   fn <- cm_tbl["LateMinor", "LateMajor"]
6   precision <- tp / (tp + fp)
7   recall    <- tp / (tp + fn)
8   f1        <- 2 * precision * recall / (precision + recall)
9   c(precision = precision, recall = recall, F1 = f1)
10 }
11 scores <- calc_measures(cm$table)
12
13 # Print a tidy summary
14 list(accuracy = acc,
15      precision = scores["precision"],
16      recall    = scores["recall"],
17      F1        = scores["F1"])
```

Majorly Late? Steps 4 - 5 Results Part 2

\$accuracy

Accuracy

0.8028169

\$precision

precision

0.2352941

\$recall

recall

0.8

\$F1

F1

0.3636364

Model Metrics (Augmented Data)

- **Accuracy** **0.80** – overall correct classifications

Model Metrics (Augmented Data)

- **Accuracy 0.80** – overall correct classifications
- **Recall 0.80** – captures most *LateMajor* cases

Model Metrics (Augmented Data)

- **Accuracy 0.80** – overall correct classifications
- **Recall 0.80** – captures most *LateMajor* cases
- **Precision 0.24** – many false alarms on *LateMajor*

Model Metrics (Augmented Data)

- **Accuracy 0.80** – overall correct classifications
- **Recall 0.80** – captures most *LateMajor* cases
- **Precision 0.24** – many false alarms on *LateMajor*
- **F1 0.36** – balance between precision & recall

Model Metrics (Augmented Data)

- **Accuracy 0.80** – overall correct classifications
- **Recall 0.80** – captures most *LateMajor* cases
- **Precision 0.24** – many false alarms on *LateMajor*
- **F1 0.36** – balance between precision & recall
- Lower accuracy is expected when you start catching the rare class.